

SOTERIA AI TECHNOLOGIES INC. · DELAWARE C-CORPORATION

Capital Efficiency & Unit Economics Memo

How Global Office converts an AI-native architecture into software margins, and what \$1.5M of pre-seed capital buys.

Global Office is operated by Soteria AI Technologies Inc., a Delaware C-Corporation. Global Office is not a productivity SaaS application; it is an AI-native platform — an intelligence layer that orchestrates business operations. That distinction drives the economics in this memo: cost to serve is dominated by metered inference rather than headcount, and value expands with intelligence consumption rather than seats alone.

1. Unit economics at a glance

Metric	Year 1	Year 3 target	Basis
Blended ARPO	\$2.4K–4.2K	\$5.6K–8.0K	Per organization per year across Starter, Growth and Scale tiers
Gross margin	76–81%	84–88%	Infrastructure plus metered AI cost, net of support
Cost to serve / org	\$14–22 / mo	\$6–9 / mo	Compute, storage, egress and model inference
CAC (blended)	\$900–1,600	\$700–1,200	Founder-led plus partner channel
CAC payback	6–9 mo	<4 mo	Gross-margin adjusted
LTV : CAC	n/a (pre-revenue)	≥ 5x	36-month horizon at target retention
Net revenue retention	—	115–135%	Seat growth, tier upgrades and AI credit expansion

All figures are forward-looking assumptions to be re-baselined against design-partner data. They are stated as ranges deliberately; no revenue has been recognized to date.

2. Why the margin structure holds

- Single multi-tenant Postgres substrate with row-level isolation — no per-customer infrastructure, so marginal orgs land on a largely fixed cost base.
- Inference is metered and tiered. Routine classification runs on small models; deliberation runs on frontier models only when an executive decision requires it. AI credits are priced with a margin floor, so heavy usage remains gross-margin accretive.
- Organizational memory reduces cost over time: repeated questions resolve against stored context and prior deliberations rather than a fresh full-context model call.
- Support cost falls with product maturity because the platform explains its own reasoning and exposes an audit trail, reducing escalations.

3. Capital efficiency: the platform is already built

The unusual feature of this round is that construction risk is largely retired. Fifteen-plus modules, the graph substrate, the cognition layer, workflow automation, meetings intelligence, predictive alerts, the trust and compliance surface, and the investor and administrative back office are shipped and running in production. Capital is therefore allocated to commercialization and evidence, not to building the first version.

Allocation	% of \$1.5M	What it converts into
Product & infrastructure	42%	Enterprise hardening, two vertical editions, scale capacity
Go-to-market	35%	30-50 design partners, first paying cohort, measured CAC
Sales, success & AI	23%	Repeatable onboarding, expansion motion, inference budget

4. Milestones the round underwrites

- Month 6 — 20 signed design partners, first paid conversions, baseline retention and activation data.
- Month 12 — measured CAC and payback on a real cohort, published NRR, two vertical editions in market.
- Month 18 — repeatable channel economics and a defensible intelligence dataset sufficient to price a seed round on evidence rather than narrative.

5. Sensitivity

The three variables that most affect outcome are inference cost per active organization, conversion rate from design partner to paid, and expansion into higher tiers. A 30% adverse move in inference pricing compresses year-three gross margin by roughly four to six points, which is absorbable. A halving of design-partner conversion is the material risk and is mitigated by pricing flexibility, local-currency plans and a partner-led channel.